

Long-form localization selects earlier in the network

bar height: share of selected heads at a top-10% budget · 19 model-task pairs, each weighted equally

long-form

mean depth 0.46

single-token

mean depth 0.53

input $\frac{1}{4}$ half-way $\frac{3}{4}$ output

Depth through the network

