



# Summarization (Long) · Long-Form Eval, Single-Token Steer · w/o R+F Filter


## Gemma 3-12B




## Qwen 1.5-14B




## Qwen 1.5-32B



## OLMo 2-13B



## Falcon3-10B



Top-K fraction of concept-sensitive attention heads