



Summarization (Long) · Single-Token Eval, Long-Form Steer · Token Matching

OLMo 2-13B



Qwen 1.5-14B



Top-K fraction of concept-sensitive attention heads