



Summarization (Single) · Long-Form Eval, Single-Token Steer · w/o R+F Filter


Gemma 3-12B




Qwen 1.5-14B




Qwen 1.5-32B



OLMo 2-13B



Falcon3-10B



Top-K fraction of concept-sensitive attention heads