



Localization: Verse (Long) · Long-Form Eval, Single-Token Steer · w/o R+F Filter


Qwen 1.5-14B



OLMo 2-13B



Qwen 1.5-32B



Top-K fraction of concept-sensitive attention heads