



Localization: Verse (Single) · Single-Token Eval, Long-Form Steer · Token Matching


Qwen 1.5-14B



OLMo 2-13B



Vicuna 13B



Top-K fraction of concept-sensitive attention heads