



Summarization (Long) · Single-Token Eval, Long-Form Steer · Token Matching


Qwen 1.5-14B




Qwen 1.5-32B



OLMo 2-13B



Gemma 3-12B



Top-K fraction of concept-sensitive attention heads